

# Protean Support Assistant

Technical Documentation & Architecture Overview

**Date:** December 11, 2025

## 1. Project Overview

The Protean Support Assistant is a **Retrieval-Augmented Generation (RAG)** chatbot designed to provide accurate, professional answers regarding GRA (Ghana Revenue Authority) E-commerce compliance. Unlike a standard chatbot that might hallucinate facts, this system is grounded in a specific dataset (your FAQ document), ensuring high accuracy while maintaining a conversational tone.

## 2. How It Works (The Workflow)

The system follows a 4-step process to answer every user query:

- **Step 1: Ingestion**

We read the source Word document (*GRA\_Ecommerce\_Chatbot\_Intent\_FAQs.docx*), extract the Questions and Answers, and store them in a structured JSON database.

- **Step 2: Vectorization (The Search Index)**

We convert every FAQ question into a mathematical vector (a list of numbers) using **TF-IDF**. This creates a 'map' of our knowledge base.

- **Step 3: Retrieval (The Search)**

When a user asks a question, we convert their text into a vector and compare it against our map using **Cosine Similarity**. This finds the most relevant FAQ entry mathematically.

- **Step 4: Generation (The AI)**

We take the *correct answer* from our database and the *user's question*, and send them to **Google Gemini AI**. The AI then rewrites the answer to be friendly, professional, and conversational.

## 3. Deep Dive: The Search Mechanism

### **Algorithm: TF-IDF (Term Frequency-Inverse Document Frequency)**

We chose TF-IDF over complex vector databases (like Pinecone or Chroma) for this specific use case.

#### **Why?**

- **Speed:** It is instant for datasets under 100,000 documents.

- **Keyword Precision:** In compliance/legal contexts, specific words (like 'VAT', 'TIN', 'Non-resident') matter more than general semantic meaning. TF-IDF is excellent at exact keyword matching.
- **Simplicity:** It runs locally without needing external servers or heavy API calls for embeddings.

## 4. Technology Stack & Rationale

Component	Technology	Why We Used It
Language	Python 3.10+	The standard for AI/ML development with vast library support.
Web Framework	Flask	Lightweight, fast, and perfect for microservices. We didn't need the overhead of a full stack framework.
AI Model	Gemini 2.5 Flash	Google's latest efficient model. It provides high intelligence with very low latency.
Search Engine	Scikit-Learn	Industry-standard library for machine learning. We use its TfidfVectorizer and CountVectorizer.
Frontend	Bootstrap 5 + JS	Ensures a professional, responsive UI that works on mobile and desktop without needing a separate framework.
Data Parsing	python-docx	Allows us to directly read the raw MS Word file provided by the client.

## 5. Key Files Explained

- **rag\_service.py:** The 'Brain'. This file contains the RAGChatbot class which handles loading data, training the search index, and talking to Gemini.
- **parse\_faqs.py:** The 'Loader'. A utility script that converts the messy Word document into clean JSON data.
- **app.py:** The 'Server'. Handles HTTP requests from the browser and connects the frontend to the backend.
- **extracted\_faqs.json:** The 'Database'. A local file acting as our knowledge base.